Agent Architectures & SDK Ecosystem

A Detailed Technical Read Document (OpenAI-Centric)

---

1. Introduction

Modern AI agents are not just language models — they are **orchestrated systems** that combine reasoning, memory, planning, tool use, and autonomous execution.

This document provides:

- A deep dive into **agent architectures**
- A breakdown of **major SDK ecosystems**
- Architectural trade-offs
- Production design considerations
- Decision guidance for enterprise systems

---

2. What Is an Agent Architecture?

An agent architecture defines:

- How decisions are made
- How tasks are decomposed
- How tools are invoked
- How state is maintained
- How feedback loops operate
- How multiple agents coordinate

It is the **structural blueprint** behind intelligence.

3. Core Architectural Patterns

3.1 Single-Agent Architecture

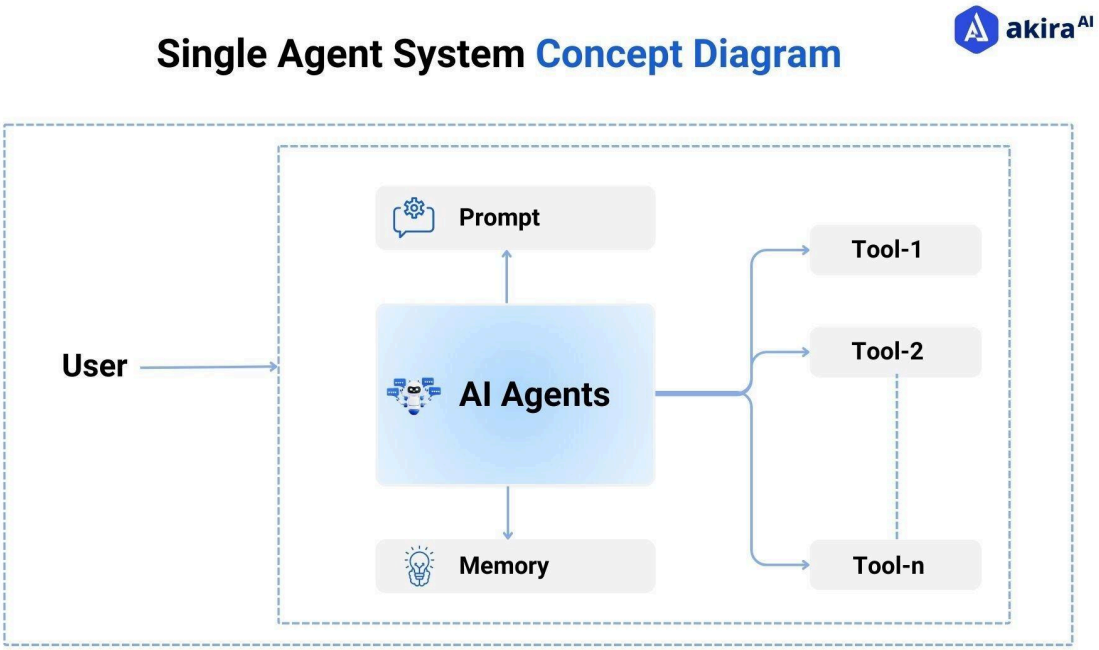

Structure

User → Agent → LLM → Tool → LLM → Output

Components

- One LLM
- Tool registry
- Memory store
- Execution loop
- State handler

How It Works

1. User provides goal
2. LLM reasons
3. Selects tool
4. Tool executes
5. Result fed back to LLM
6. Loop continues until goal achieved

Advantages

- Simple
- Cost-efficient
- Easier debugging
- Lower latency

Limitations

- Cognitive overload
- Less modular
- Harder to scale for complex tasks

3.2 ReAct (Reason + Act)

ReAct combines reasoning and action in iterative loops.

Thought → Action → Observation → Thought → Action

This pattern:

- Improves reliability
- Reduces hallucination
- Makes tool usage explicit

It is foundational to most modern OpenAI-based agent designs.

3.3 Planner–Executor Architecture

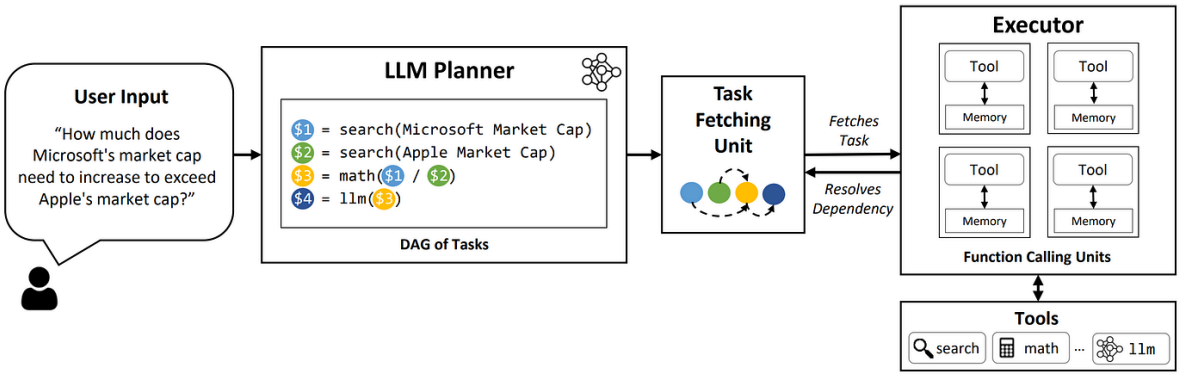

Structure

Two logical roles:

- Planner
- Executor

Flow

Goal ⟶ Planner ⟶ Task List

↓

Executor ⟶ Tools

Why It Matters

Separates:

- Strategic reasoning (Planner)
- Tactical execution (Executor)

Benefits

- Reduced reasoning load
- Parallel task execution
- Improved monitoring
- Better observability

Enterprise Use Cases

- Data analytics workflows
- DevOps automation
- Research pipelines
- Multi-step business processes

3.4 Multi-Agent Architecture

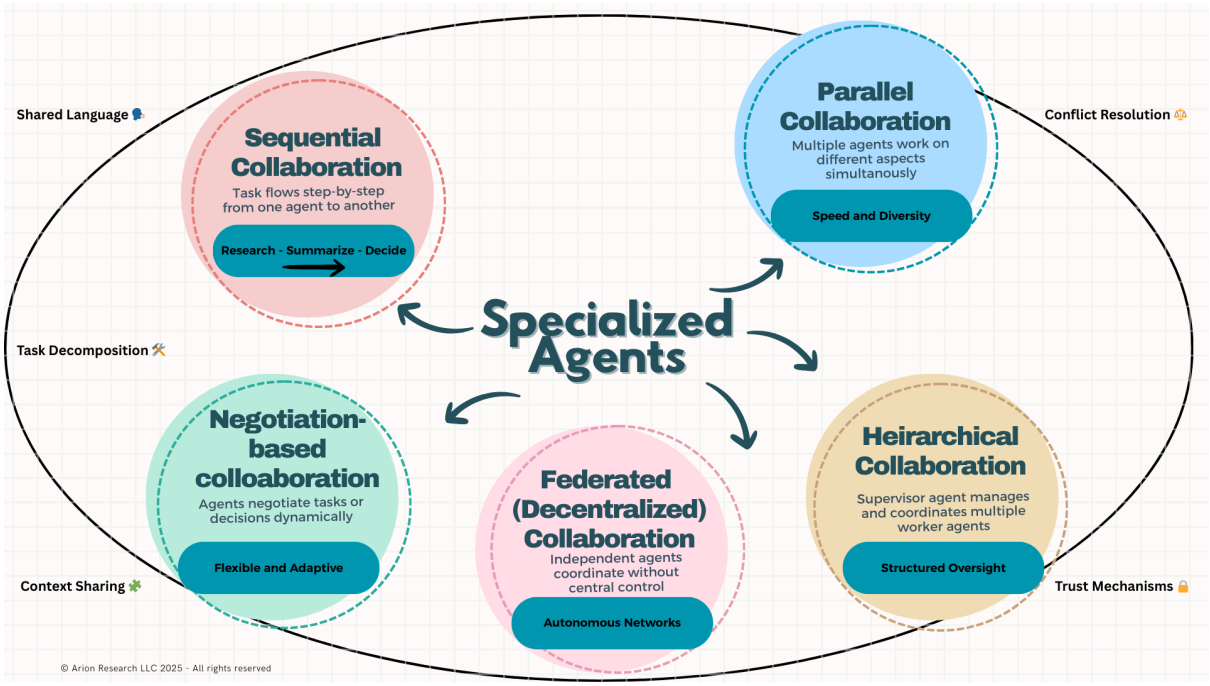

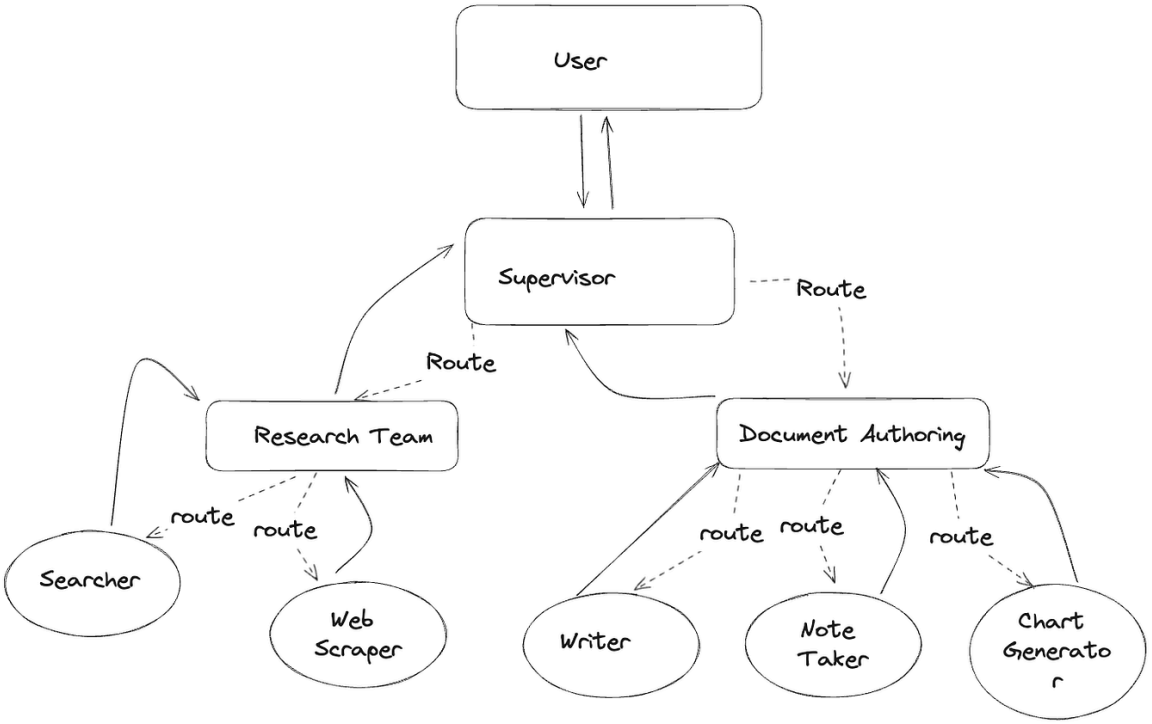

Instead of one agent doing everything:

Create specialized agents:

- Research Agent
- Planning Agent
- Coding Agent
- Validation Agent
- Memory Agent

Coordination Models

- Hierarchical (Manager → Workers)
- Peer-based
- Swarm-based
- Graph-directed workflows

Advantages

- Specialization improves quality
- Parallel execution
- Modular scaling
- Clear responsibility separation

Trade-offs

- Higher cost
- Complex orchestration
- Observability challenges

3.5 Graph-Based Architectures

Agents represented as state machines or directed graphs.

Nodes:

- LLM reasoning
- Tool execution
- Validation
- Memory update

Edges:

- Conditional routing
- Retry paths
- Error handling

This provides:

- Deterministic control
- High reliability
- Production-grade stability

Common in enterprise AI orchestration.

4. SDK Ecosystem for Agent Development

Architecture is conceptual.

SDKs are implementation tools.

4.1 OpenAI Native SDK

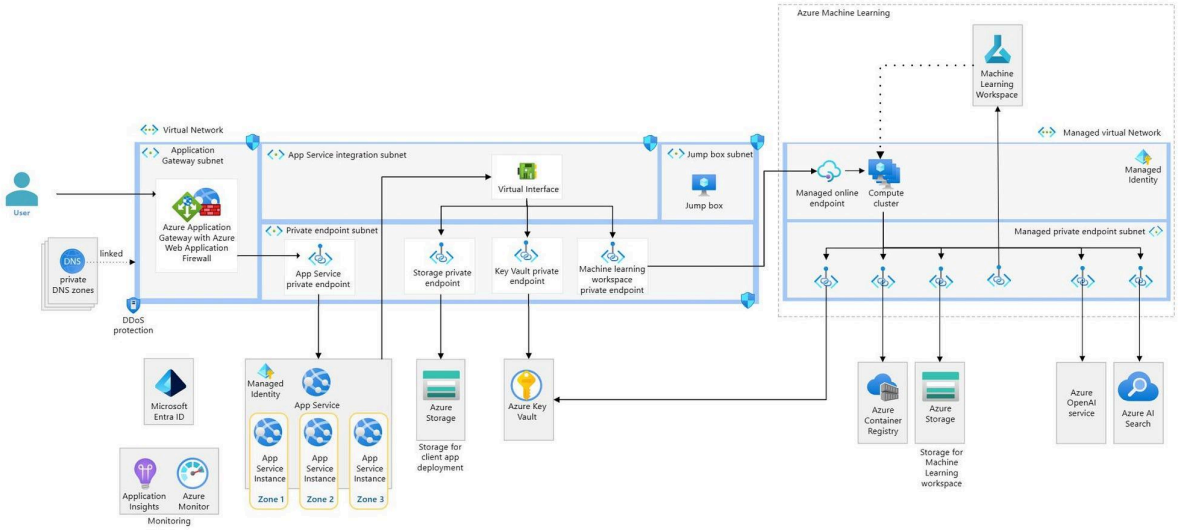

Core Capabilities

- Tool calling
- Structured JSON outputs
- Embeddings API
- Streaming
- Multimodal inputs
- System prompts & role separation

Why Use Native SDK?

- Maximum control
- Production stability
- Custom orchestration
- Better cost control
- Lower abstraction overhead

Best For

- Enterprise systems
- High reliability requirements
- Custom control loops
- Regulated environments

4.2 LangChain

Framework for rapid LLM application development.

Provides

- Agent abstractions
- Tool wrappers
- Memory integrations
- RAG components
- Prompt templates

Strengths

- Rapid prototyping
- Large ecosystem
- Many integrations

Weaknesses

- Abstraction complexity
- Debugging challenges
- Version instability

Best for:

- Proof-of-concepts
- Research environments
- Fast iteration

4.3 LangGraph (LangChain Extension)

Graph-based orchestration system.

Features

- Stateful graph control
- Retry policies
- Conditional routing
- Multi-agent orchestration

Ideal for:

- Production multi-agent systems
- Deterministic workflows

4.4 Microsoft AutoGen

Multi-agent conversational framework.

Focus

- Agent-to-agent dialogue
- Role-based collaboration
- Research experimentation

Best for:

- Academic research
- Experimental systems
- Agent collaboration modeling

4.5 Microsoft Semantic Kernel

Enterprise-focused framework.

Features

- Plugin architecture
- Planning abstractions
- Strong .NET integration
- Enterprise readiness

Best for:

- Microsoft-heavy enterprise stacks

4.6 CrewAI

Lightweight multi-agent orchestration system.

Focus

- Role-based agent teams
- Simple coordination
- Rapid experimentation

Good for:

- Multi-agent prototyping

---

5. Architecture vs SDK (Critical Distinction)

Concept	Meaning
Architecture	Design pattern of intelligence
SDK	Implementation toolkit

Example:
You can implement a Planner–Executor model using:

- OpenAI SDK directly
- LangGraph
- Semantic Kernel
- Custom orchestration

SDK does not define intelligence strategy — architecture does.

---

6. Production Design Considerations

6.1 Observability

You must monitor:

- Token usage
- Tool call frequency
- Failure rate
- Retry loops
- Latency
- Cost per task

6.2 Guardrails

Include:

- Output validation
- Schema enforcement
- Safety filtering
- Tool permission checks
- Rate limits

6.3 Cost Modeling

Multi-agent systems:

- Multiply token usage
- Increase tool calls
- Add latency

Planner–Executor reduces waste by isolating reasoning layers.

6.4 Failure Modes

Common problems:

- Infinite loops
- Tool misuse
- Hallucinated tool parameters
- Context overflow
- Memory corruption
- Cascading reasoning errors

Graph-based or controlled loops mitigate this.

---

7. Decision Guide

Requirement	Suggested Approach
Simple automation	Single-Agent
Multi-step workflows	Planner–Executor
High specialization	Multi-Agent
Strict control	Graph-based
Enterprise reliability	Native SDK + custom orchestration
Rapid prototyping	LangChain
Research experimentation	AutoGen

---

8. Recommended Enterprise Pattern

For production systems:

- Use OpenAI native APIs
- Build custom orchestration layer
- Implement Planner–Executor structure
- Add memory via embeddings
- Add validation layer
- Add observability stack
- Add cost monitoring
- Implement retry policies

Frameworks accelerate development — but production often requires controlled orchestration.

---

9. Final Summary

Modern AI agent systems consist of:

- Structured architecture
- Orchestration strategy
- Tool integration
- Memory systems
- Feedback loops
- Observability
- Guardrails

Choosing the right architecture and SDK depends on:

- Complexity
- Scale
- Reliability needs
- Budget
- Ecosystem compatibility